

## Assignment #2: Greedy Algorithms Solutions

1. (10 marks) Consider a communications graph. Each edge of the connected graph  $G = (V, E)$  represents a communication link between sites (represented as nodes). Each edge  $e$  has a bandwidth  $b_e$ .

Each pair of nodes  $u, v \in V$  needs to be able to communicate. For any  $u, v$  - path  $P$  the *bottleneck transmission rate*  $b(P)$  of  $P$  is the minimum bandwidth of any edge it contains. In other words,  $b(P) = \min_{e \in P} b_e$ . The *best achievable bottleneck rate* for a pair  $u, v \in V$  is the maximum, over all  $u - v$  paths  $P$  in  $G$ , of the value  $b(P)$ . Our goal is to determine a set of  $u - v$  paths for each pair  $u, v \in V$  with best achievable bottleneck rate.

Fortunately, we can construct a spanning tree  $T$  of  $G$  such that for every pair of nodes  $u, v \in V$  the unique  $u - v$  path in the tree  $T$  actually achieves the best achievable bottleneck rate for  $u, v$  in  $G$ .

Give an efficient algorithm to construct such a spanning tree. Your algorithm should construct a spanning tree  $T$  in which, for each  $u, v \in V$ , the bottleneck rate of the  $u - v$  path in  $T$  is equal to the best achievable bottleneck rate for the pair  $u, v$  in  $G$ . Prove the correctness of your algorithm.

### Sample Solutions.

**Claim.** A minimum spanning tree, computed with each edge cost equal to the negative of its bandwidth is a spanning tree such that for each  $u, v \in V$ , the bottleneck rate of the  $u - v$  path in  $T$  is equal to the best achievable bottleneck rate for the pair  $u, v$  in  $G$ .

**Proof.** Let's first assume all edge weights are unique. Assume for a contradiction that the claim is not true. Then there exists some pair for nodes  $u, v$  for which the path  $P$  in the minimum spanning tree does not have a bottleneck rate as high as some other  $u - v$  path  $P'$ . Let  $e = (x, y)$  be an edge of minimum bandwidth on the path  $P$ ; note that  $e \notin P'$ . Furthermore,  $e$  has the smallest bandwidth of any edge in  $P \cup P'$ . Now, using the edges in  $P \cup P'$  other than  $e$  it is possible to travel from  $x$  to  $y$  (for example, by going from  $x$  back to  $u$  via  $P$ , then to  $v$  via  $P'$ , then to  $y$  via  $P$ ). Thus, there is a simple path from  $x$  to  $y$  using these edges and so there is a cycle  $C$  on which  $e$  has the minimum bandwidth.

This means that in our minimum spanning tree instance,  $e$  has the max. cost on the cycle  $C$ ; but then the minimum spanning tree algorithm would not have included  $e$  but rather some other edge on the cycle  $C$  contradicting that we have a minimum spanning tree.

Now, if the edge costs are not all distinct we need to find a way to slightly alter the edge weights without compromising the solution. Consider perturbing all edge bandwidths by extremely small amounts so they become distinct, and then defining a minimum spanning tree. We therefore refer, for each edge  $e$  to a *real bandwidth*  $b_e$  and a *perturbed bandwidth*  $b'_e$ . In particular, we choose perturbations small enough so that if  $b_e > b_f$  for edges  $e$  and  $f$ , then also  $b'_e > b'_f$ . Our tree has the best bottleneck rate for all pairs, under the perturbed bandwidths. But suppose that the  $u - v$  path  $P$  in this tree did not have the best bottleneck rate if we consider the original, real bandwidths; say there is a better path  $P'$ . Then there is an edge  $e$  on  $P$  for which  $b_e > b_f$  for all edges  $f$  on  $P'$ . But the perturbations were so small that they did not cause any edges with distinct bandwidths to change the relative order of their bandwidth values, so it would follow that  $b'_e > b'_f$  for all edges  $f$  on  $P'$ , contradicting our conclusion that  $P$  had the best bottleneck rate with respect to the perturbed bandwidths.

**Complexity.** Same as Prim's or Dijkstra's or  $\mathcal{O}(n + m \log n)$ .

2. (10 marks) Given a set  $P$  of  $n$  people. Suppose the  $i^{\text{th}}$  person claims to know  $d_i$  other people in  $P$ . Determine in polynomial time if  $P$  is a feasible set by constructing a greedy algorithm. By feasible, we mean that it's possible for each person to know the number of people they claim to know. Prove that

your algorithm correctly determines if such a set up is possible. HINT: Think of the  $i^{th}$  person being represented by a vertex  $v_i$  and the number of people that the person knows as the degree  $d_i$  of  $v_i$ .

### Sample Solutions

Clearly if any of the degrees  $d_i$  is equal to 0, then this must be an isolated node in the graph, so we can delete  $d_i$  from the list and continue by recursion on the smaller instance.

Otherwise, all  $d_i$  are positive. We sort the numbers, relabeling as necessary so that  $d_1 \geq d_2 \geq \dots \geq d_n$ . We now consider the list of numbers:

$$L = \{d_1 - 1, d_2 - 1, \dots, d_{d_n} - 1, d_{d_n+1}, \dots, d_{n-2}, d_{n-1}\}$$

So we subtract 1 from each of the first  $d_n$  numbers and drop the last number. We claim that

*“there exists a graph whose degrees are equal to the list  $d_1, d_2, \dots, d_n$  if and only if there exists a graph whose degrees form the list  $L$ . ”*

Assuming this claim, we can proceed recursively.

### Proof of claim.

First if there is a graph with degree sequence  $L$ , then we can add an  $n^{th}$  node with neighbours equal to the nodes  $v_1, v_2, \dots, v_{d_n}$  thereby obtaining a graph with degree sequence  $d_1, \dots, d_n$ . Conversely, suppose there is graph with degree sequence  $d_1, d_2, \dots, d_n$  where we have  $d_1 \geq d_2 \geq \dots \geq d_n$ . We must show that in this case, there is in fact such a graph where node  $v_n$  is joined to precisely the nodes  $v_1, v_2, \dots, v_{d_n}$ . This will allow us to delete node  $n$  and obtain list  $L$ .

Consider any graph  $G$  with degree sequence  $d_1, \dots, d_n$ ; we show how to transform  $G$  into a graph where  $v_n$  is joined to  $v_1, v_2, \dots, v_{d_n}$ . If this property does not already hold, then there exists  $i < j$  so that  $v_n$  is joined to  $v_j$  but not  $v_i$ . Since  $d_i \geq d_j$ , it follows that there must be some  $v_k$  not equal to any of  $v_i, v_j, v_n$  with the property that  $(v_i, v_k)$  is an edge but  $(v_j, v_k)$  is not. We now replace these two edges by  $(v_i, v_n)$  and  $(v_j, v_k)$ . This keeps all degrees the same; and repeating this transformation will convert  $G$  into a graph with the desired property.

### Complexity

We take  $\Theta(n \log n)$  time to sort the  $n$  degrees. For each vertex  $v_i$  we need to do update  $d_i$  vertices degrees. This works out to  $\Theta(m)$  time for a total of  $\Theta(n \log n + m)$  time.